

CMSC 691

High Performance Distributed Systems

Introduction to CUDA

Dr. Alberto Cano
Assistant Professor
Department of Computer Science
acano@vcu.edu

The CPU vs the GPU purposes

CPU

- General-purpose computation
- Small number of highly specialized complex cores
- Fast execution of a single stream of instructions
- Pipelining, caching, branch prediction, out-of-order execution, interruptions
- Main DDR memory ~ 20 GB/s

GPU

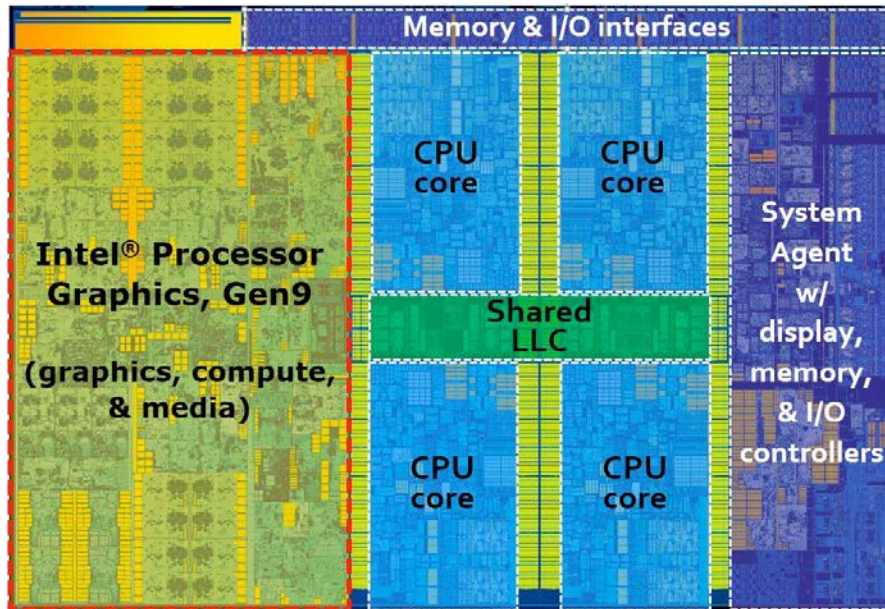
- Originally for graphics computing
- Large number of simple cores for parallel SIMD computation
- Large FP bandwidth for massive parallel number of operations
- No fancy hardware tricks except asynchronous data / execution
- GPU memory GDDR > 200 GB/s
- Future: HBM2 > 1 TB/s

CMSC 691 High Performance Distributed Systems

Introduction to CUDA

The CPU vs the GPU architectures


CPU (Skylake)



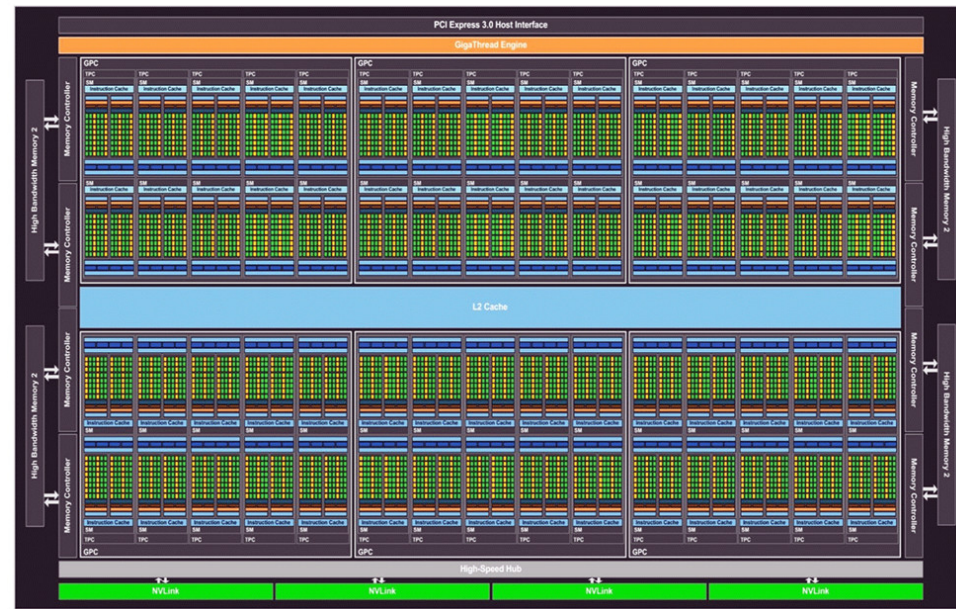
4 cores

1.75 billion transistors

110 Gflops FP32




GPU (P100)



3840 cores

15 billion transistors


10 Tflops FP32




CMSC 691 High Performance Distributed Systems

Introduction to CUDA


The GPU evolution, performance and bandwidth



Theoretical GFLOP/s

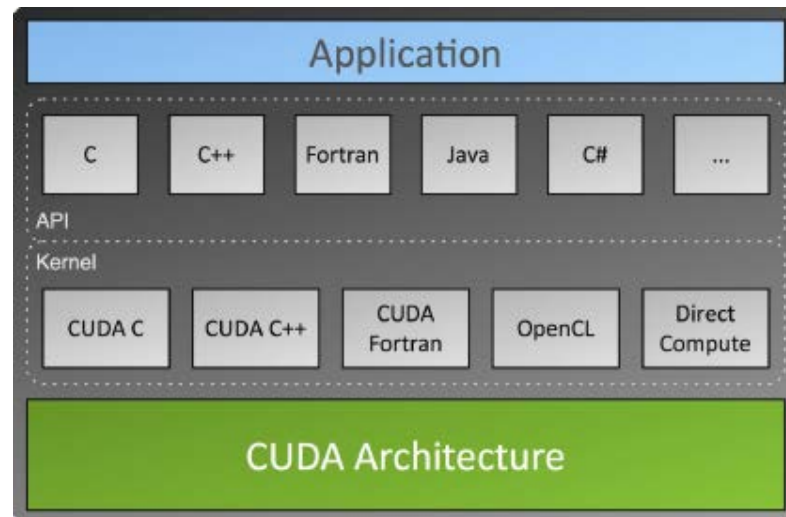


Theoretical GB/s



Exploiting the GPU

- CUDA C/C++
- CUDA Fortran
- OpenCL
- PyCUDA
- jCUDA
- Matlab
- Tensor flow
- OpenCV



Many Different Approaches

- Application level integration
- High level, implicit parallel languages
- Abstraction layers & API wrappers
- High level, explicit language integration
- Low level device APIs

Some interesting libraries for high-performance computing



- Thrust: parallel algorithms and data structures, e.g. transformation, reductions, sorting




- cuDNN: CUDA Deep Neural Network is a GPU-accelerated library of primitives for neural networks



- nvGRAPH: Graph Analytics Library for GPUs

- cuRAND: Random Number Generation library



- cuBLAS: Basic Linear Algebra Subroutines


Terminology

- Host: the CPU and its memory, typically *h_variable*
- Device: the GPU and its memory, typically *d_variable*

Thread Hierarchy

- Thread: smallest sequence of programmed instructions
- Warp: group of 32 threads scheduled/executed in a multi-processor
- Block: group of threads defined by the programmer
- Grid: group of blocks defined by the programmer
- Multi-processor: group of physical CUDA cores

CUDA threads are extremely lightweight, almost no creation overhead, context switching is essentially free



Memory Hierarchy

- Thread local memory

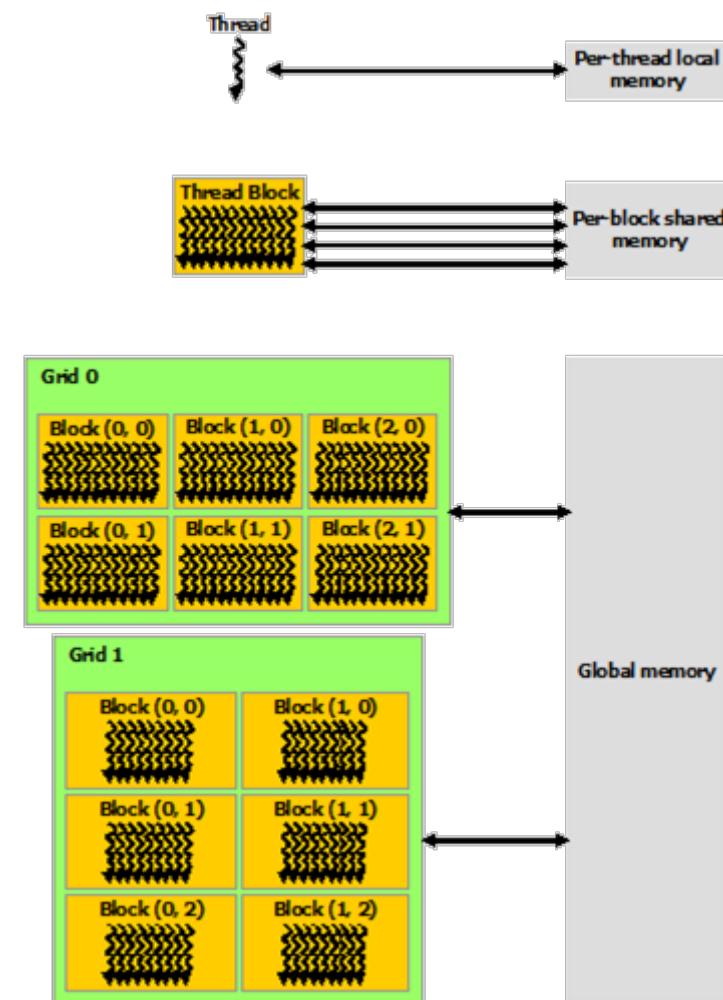
Each thread has private local memory

- Block shared memory

Each thread block has shared memory visible to all threads of the block and with the same lifetime as the block

- Device global memory

All threads have access to the same global memory. Lifetime of the program

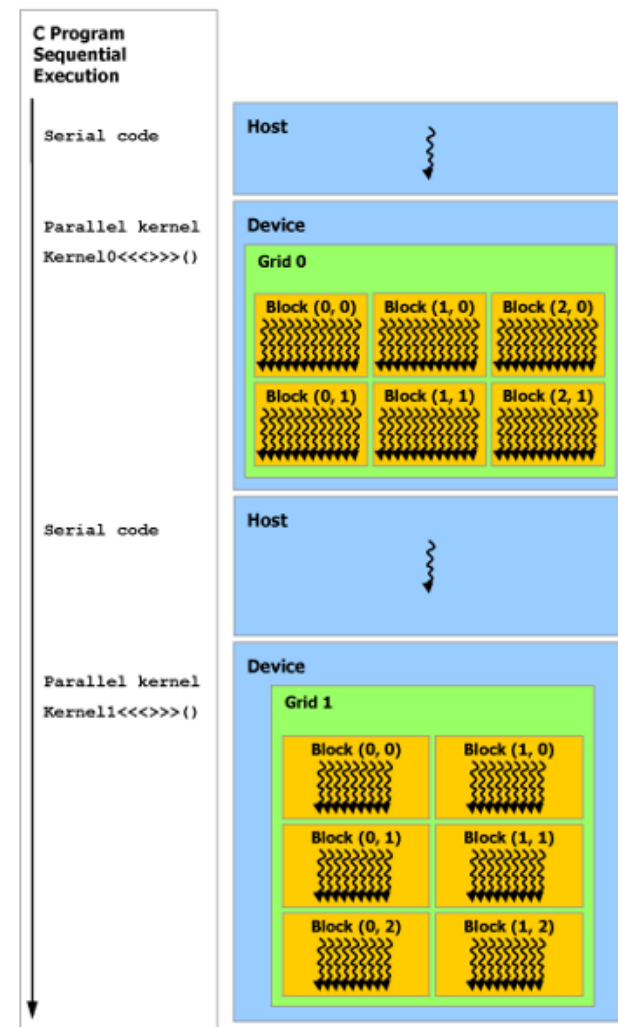


CMSC 691 High Performance Distributed Systems

Introduction to CUDA

Heterogeneous Programming

- Programs interleave code executed in the host (CPU) and GPU (device)
- GPU code is written in *kernel* functions
- Typically, the sequence of a program is:
 1. Allocate GPU memory for inputs
 2. Allocate GPU/CPU memory outputs
 3. Copy inputs from host to device
 4. Execute GPU code
 5. Copy outputs from device to host



Structure of a CUDA program

```
__global__ void vectorAdd(float* a, float* b, float* c) { // GPU kernel
    c[tid] = a[tid] + b[tid];
}

void main(void) {
    // Allocate CPU and GPU memory      (d_a, d_b, d_c, h_a, h_b, h_c)
    // Copy host data to device memory   h_a -> d_a, h_b -> d_b
    // Execute GPU kernel
    deviceKernel <<< 1,size >>> (d_a, d_b, d_c);
    // Copy device results to host memory d_c -> h_c
}
```

- NVIDIA compiler `nvcc` (can be used for programs with no GPU code)

\$ nvcc -o myprogram mycode.cu

- Threads are grouped into multi-dimensional thread blocks
- Thread blocks are grouped into a multi-dimensional grid

Thread indexing

- Thread space 3D grid of 3D blocks
- Built-in variables

gridDim.x gridDim.y gridDim.z



blockDim.x blockDim.y blockDim.z

blockIdx.x blockIdx.y blockIdx.z

threadIdx.x threadIdx.y threadIdx.z

- Indexing 1D grid of 1D blocks:

*int tid = blockIdx.x * blockDim.x + threadIdx.x;*



Memory allocation and transfer

```
cudaMallocHost(&h_ptr, count * sizeof(datatype));  
cudaMalloc(&d_ptr, count * sizeof(datatype));  
  
cudaMemcpy(dst_ptr, src_ptr, count * sizeof(datatype), cudaMemcpyKind);  
  
cudaFreeHost(h_ptr);  
cudaFree(d_ptr);
```

CUDA memory copy types

cudaMemcpyHostToDevice
cudaMemcpyDeviceToHost

Host -> Device
Device -> Host

Kernel setup

```
dim2 gridDim(1,1); // 1D grid  
dim3 blockDim(256,1,1); // 1D block with 256 threads  
  
MyKernel <<< gridDim, blockDim >>> (d_data);
```

VectorAdd example

```
__global__ void vectorAdd(float *A, float *B, float *C) {
    int tid = blockDim.x * blockIdx.x + threadIdx.x;
    C[tid] = A[tid] + B[tid];
}

void main(void)
{
    float *h_A = (float *)malloc(numElements * sizeof(float));
    ...
    cudaMalloc(&d_A, numElements * sizeof(float));
    ...
    cudaMemcpy(d_A, h_A, numElements*sizeof(float), cudaMemcpyHostToDevice);
    ...
    vectorAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A, d_B, d_C);
    cudaMemcpy(h_C, d_C, numElements*sizeof(float), cudaMemcpyDeviceToHost);
}
```

- Careful! Let's see the full working code
- Use *cuda-memcheck yourprogram* to find memory errors

CMSC 691

High Performance Distributed Systems

Introduction to CUDA

Dr. Alberto Cano
Assistant Professor
Department of Computer Science
acano@vcu.edu