

# CMSC 691


## High Performance Distributed Systems

### CUDA Asynchronous Concurrent Execution

Dr. Alberto Cano  
Assistant Professor  
Department of Computer Science  
[acano@vcu.edu](mailto:acano@vcu.edu)


GPU processing flow so far

1. Copy input data from CPU memory to GPU memory
2. Launch a GPU kernel
3. Copy results from GPU memory to CPU memory



- Kernel needs to wait input data to be transferred
- Results cannot be copied back until kernel finished
- What if we launch multiple kernels?
- Need to wait until all input data is copied to start the kernel?
- Need to wait until kernel finishes to start copying results?

## Execution of multiple kernels with no dependencies



```

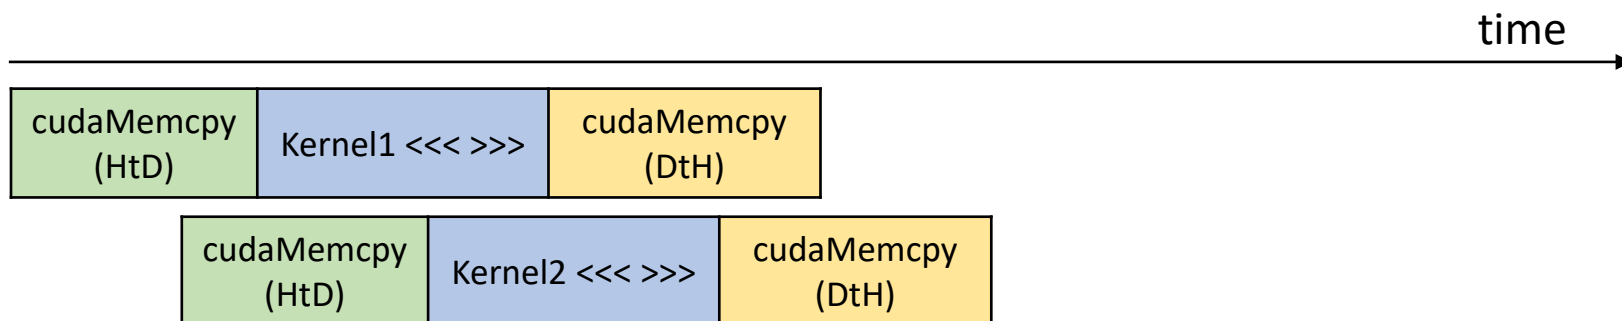
cudaMemcpy(d_input1, h_input1, size, cudaMemcpyHostToDevice);
Kernel1<<<blocks1, threads1>>>(d_input1, d_output1);
cudaMemcpy(h_output1, d_output1, size, cudaMemcpyDeviceToHost);

cudaMemcpy(d_input2, h_input2, size, cudaMemcpyHostToDevice);
Kernel2<<<blocks2, threads2>>>(d_input2, d_output2);
cudaMemcpy(h_output2, d_output2, size, cudaMemcpyDeviceToHost);

```

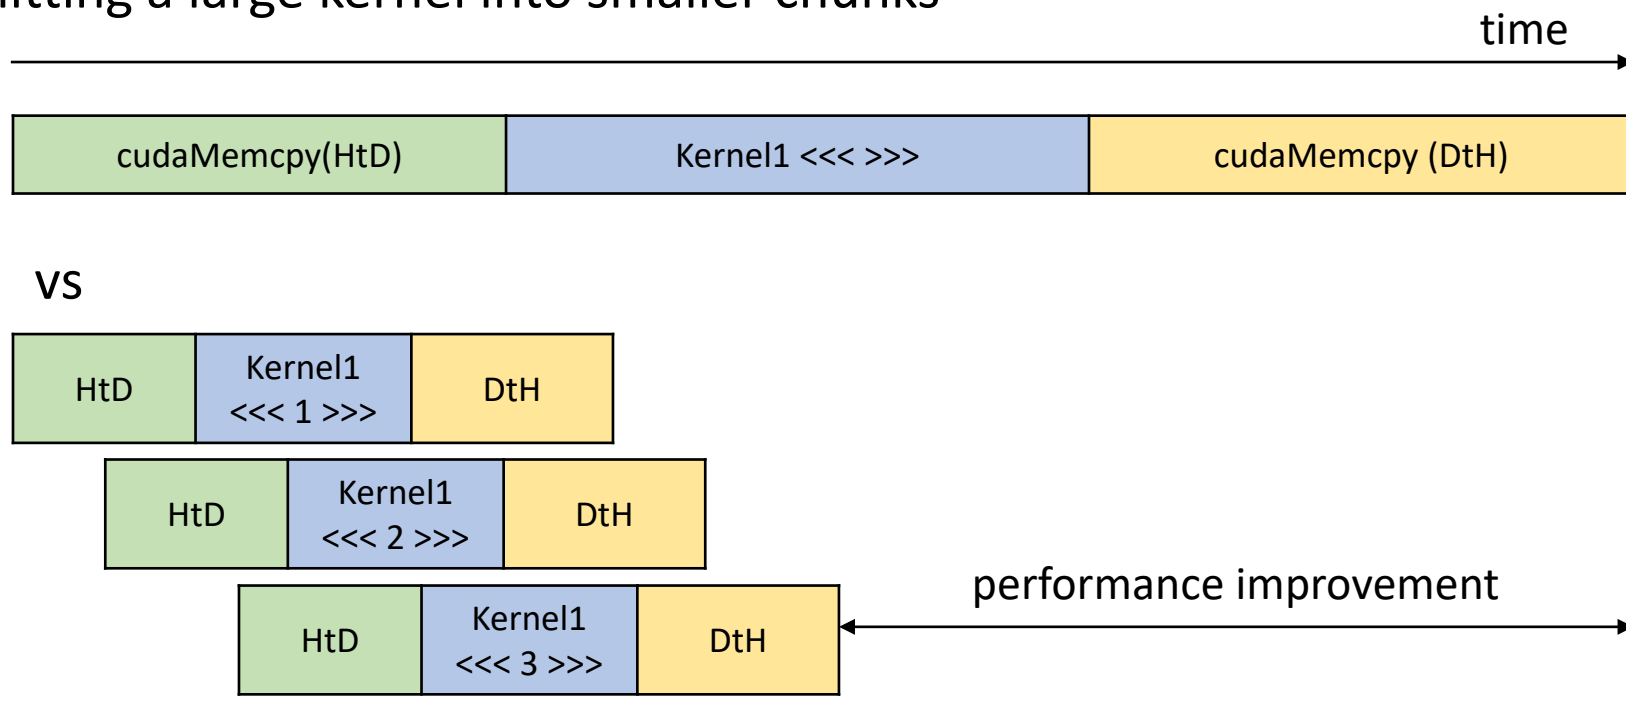
- Assuming Kernel1 and Kernel2 are data-independent
- Is it necessary to serialize the kernel executions?
- Is it necessary to serialize the data transfers?
- Can we overlap them?

## Pipelining using streams



- Independent data transfer and kernel execution flows
- Maximize the occupancy of the GPU resources
- Minimize the latency of the program
- Overlapping of data transfer and execution using CUDA streams
- Streams simulate multiple pipelines

## Splitting a large kernel into smaller chunks



- A given kernel is executed on a large data
- Divide into smaller chunks
- Multiple concurrent streams to process each chunk
- Overlapping of data transfer (HtD and DtH) and execution

## CUDA streams

- A stream is a queue of device work
- The host places work in the queue and continues on immediately
- The device schedules work from streams when resources are free
- Operations within a stream are ordered (FIFO) and cannot overlap
- Operations in different streams are unordered and can overlap

## Declaration and allocation

```
cudaStream_t stream; // single stream
```

```
cudaStreamCreate(&stream);
```

```
cudaStream_t *streams = (cudaStream_t*)malloc(n * sizeof(cudaStream_t));
```

```
for(int i = 0; i < n; i++) // multiple streams
```

```
    cudaStreamCreate(&streams[i]);
```

Scheduling a kernel to execute in a stream

- Kernels parameters:
  - Blocks setup (1D, 2D, 3D blocks)
  - Threads setup (1D, 2D, 3D threads)
  - Shared memory amount
  - Stream

*Kernel <<< blocks, threads, smem, stream >>> ();*

Default stream

- Default stream used when no stream is specified, referred as 0
- Completely synchronous w.r.t. host and device

## Kernel concurrency

- Assume *kernel* only utilizes 50% of the GPU
- Default stream

```
kernel <<<blocks,threads>>>();
```

```
kernel <<<blocks,threads>>>();
```



- Default & user streams

```
cudaStream_t stream1;
```

```
cudaStreamCreate(&stream1);
```

```
kernel <<<blocks,threads>>>();
```

```
kernel <<<blocks,threads,0,stream1>>>();
```

```
cudaStreamDestroy(stream1);
```






## Kernel concurrency

- Assume *kernel* only utilizes 50% of the GPU
- User streams

```

cudaStream_t stream1, stream2;
cudaStreamCreate(&stream1);
cudaStreamCreate(&stream2);
kernel <<<blocks,threads,0,stream1>>>();
kernel <<<blocks,threads,0,stream2>>>();
cudaStreamDestroy(stream1);
cudaStreamDestroy(stream2);

```



## Requirements for concurrency

- CUDA operations must be in different streams
- *cudaMemcpyAsync* with host from pinned/page-locked memory
  - Allocate with *cudaMallocHost()* or *cudaHostAlloc()*
  - Release with *cudaFreeHost()*
- Sufficient resources must be available
  - *cudaMemcpyAsyncs* in different directions
  - Device resources (multiprocessors, registers, blocks, etc)
- Careful *cudaMemcpyAsync* are **non-blocking!**

## Asynchronous memory transfers

- Transfers must be in a non-default stream
- Must use async memcpy
- 1 transfer per direction at a time
- Memory on the host must be pinned

```
cudaMallocHost(&h_ptr, bytes);
```

```
cudaMalloc(&d_ptr, bytes);
```

```
...
```

```
cudaMemcpyAsync(d_ptr, h_ptr, bytes, cudaMemcpyHostToDevice, &stream);
```

```
kernel << blocks, threads, smem, stream >>> (d_ptr);
```

```
cudaMemcpyAsync(h_ptr, d_ptr, bytes, cudaMemcpyDeviceToHost, &stream);
```

```
...
```

## Concurrency examples

- Synchronous

```
cudaMemcpy(...);  
kernel <<<...>>>();
```




- Asynchronous Same Stream

```
cudaMemcpyAsync(...,stream1);  
kernel <<<...,stream1>>>();
```



- Asynchronous Different Streams

```
cudaMemcpyAsync(...,stream1);  
kernel <<<...,stream2>>>();
```



## Synchronization

- Synchronize everything  
*cudaDeviceSynchronize()*
  - Blocks until all issued CUDA calls are complete (all streams)
- Synchronize host w.r.t. a specific stream  
*cudaStreamSynchronize(stream)*
  - Blocks until all issued CUDA calls in stream are complete
- Synchronize host or devices using events
  - We already used events to measure elapsed time

# CMSC 691

## High Performance Distributed Systems

### CUDA Asynchronous Concurrent Execution

Dr. Alberto Cano  
Assistant Professor  
Department of Computer Science  
[acano@vcu.edu](mailto:acano@vcu.edu)