

CMSC 691

High Performance Distributed Systems

Apache Hadoop

Dr. Alberto Cano
Assistant Professor
Department of Computer Science
acano@vcu.edu

Apache Hadoop

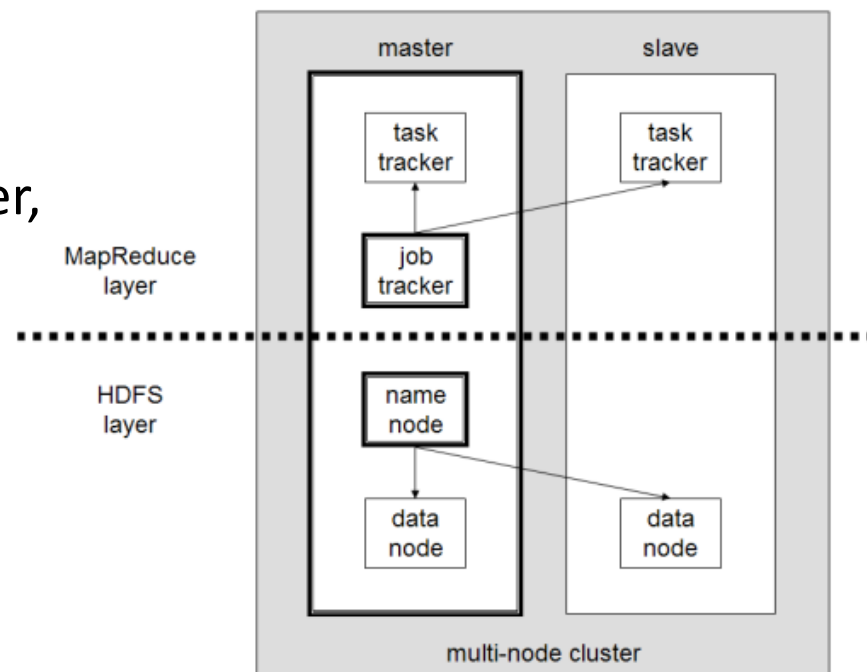
- Open-source framework for distributed storage and computing of big data sets on clusters
- Hadoop Distributed File System (HDFS): fault-tolerant, high-bandwidth, high availability distributed storage
- MapReduce: distributed big data processing infrastructure (abstract/paradigm, fault-tolerant, schedule, execution)
- Data locality: processing data local to each compute node i.e. “don't move data to workers, move workers to the data”
- Assumptions: commodity hardware is inexpensive but fail all the time, moderate number of huge write-once-read-only data files

CMSC 691 High Performance Distributed Systems

Apache Hadoop

Apache Hadoop architecture

- Hadoop Common: OS level abstractions
- Hadoop Distributed File System (HDFS)
- MapReduce engine
- Master: Job Tracker, Task Tracker, NameNode, and DataNode
- Slave/Workers: DataNode and Task Tracker

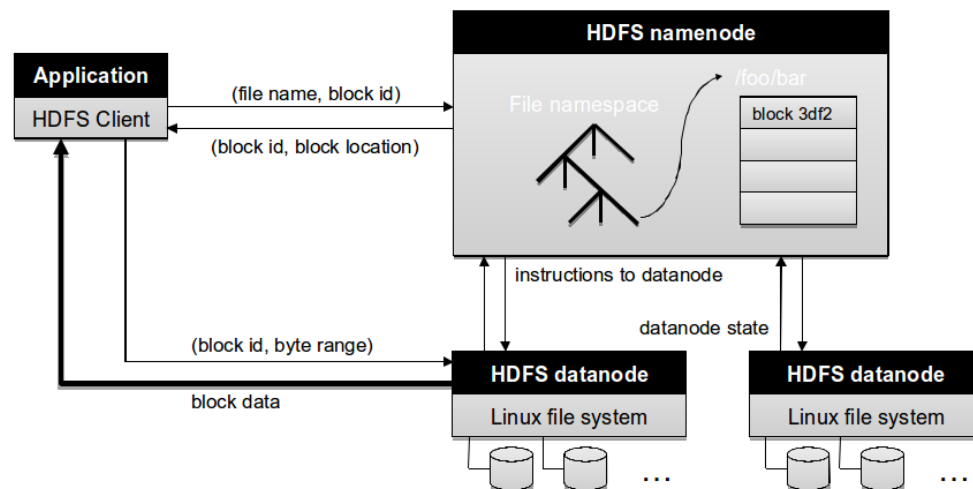


CMSC 691 High Performance Distributed Systems

Apache Hadoop

HDFS

- Single namespace for entire cluster
- NameNode: Maps a file to a file-id and list of DataNodes
- DataNode: Maps a block-id to a physical location on disk
- Files are broken up into blocks of 64 MB
- Data coherency (write-once-read-many)
- Data replication 3



HDFS NameNode

- Managing the file system namespace:
 - Holds file/directory structure, metadata, file-to-block mapping, access permissions, etc.
- Coordinating file operations:
 - Directs clients to DataNode for reads and writes
 - No data is moved through the NameNode
- Maintaining overall health:
 - Periodic communication with the DataNode
 - Block re-replication and rebalancing

HDFS DataNode

- A Block Server
 - Stores data in the local file system
 - Stores meta-data of a block (checksums)
 - Serves data and meta-data to clients
- Block Report
 - Periodically sends a report of all existing blocks to the NameNode
 - Facilitates Pipelining of Data
- Forwards data to other specified DataNodes

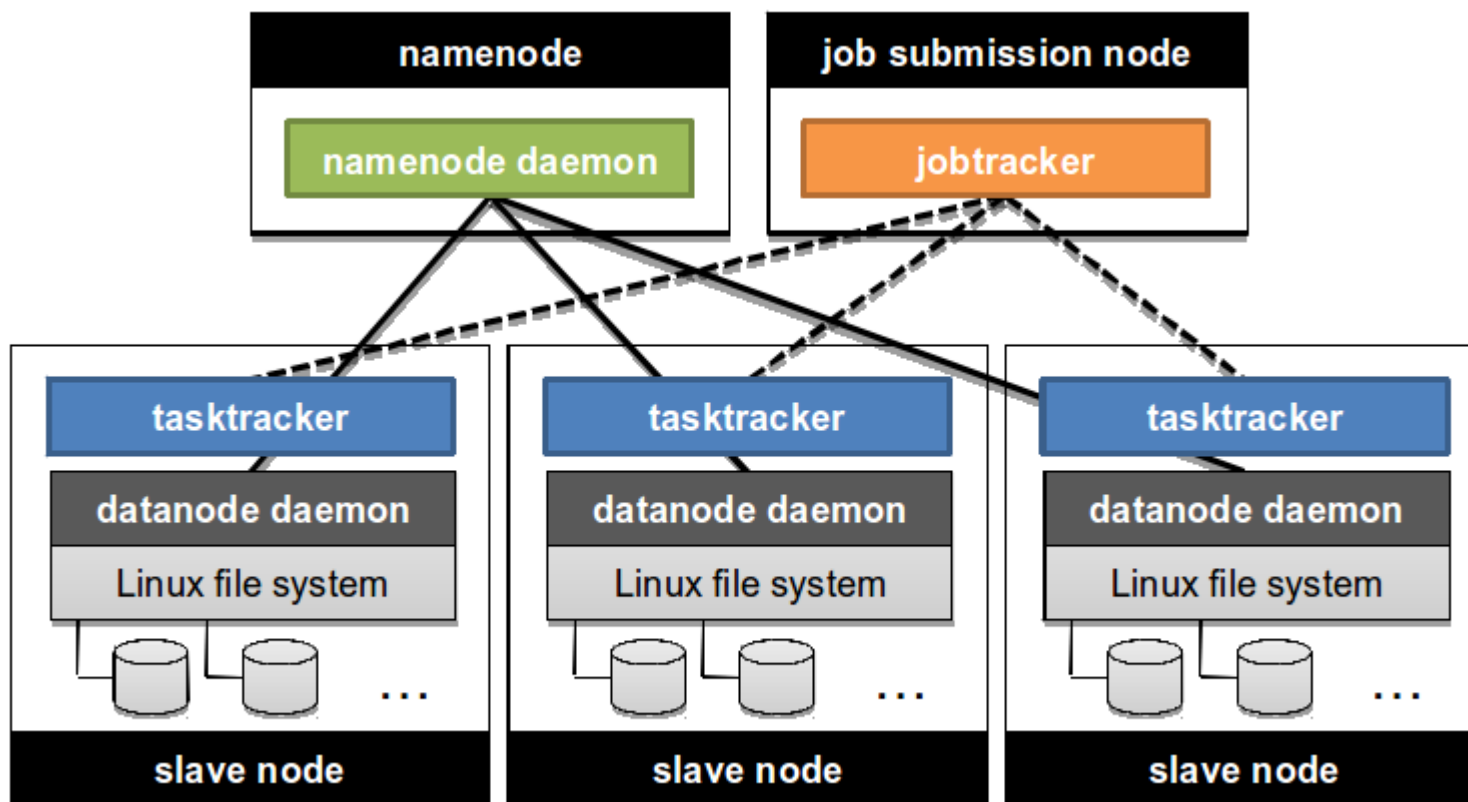
HDFS block placement

- Current strategy: default replication 3
 - One replica on one node in the local rack
 - Second replica on one node in a remote rack
 - Third replica on another node in the same remote rack
 - Additional replicas are randomly placed
- Clients read from nearest replica
- Data correctness validated through checksums
- Fault-tolerance: NameNode is a single point of failure!

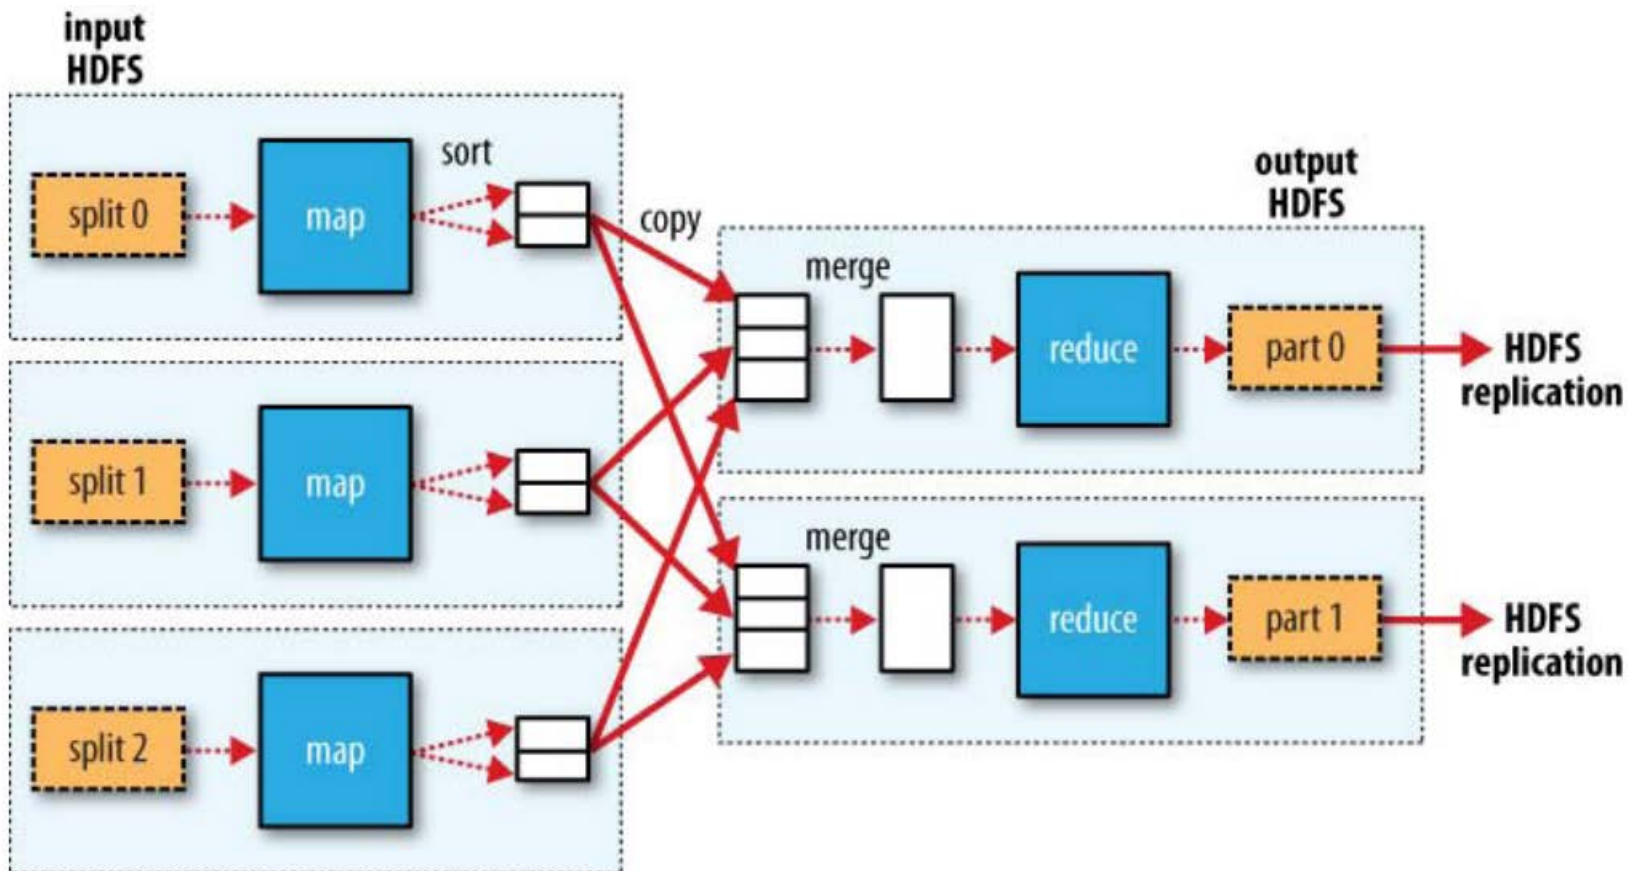
CMSC 691 High Performance Distributed Systems

Apache Hadoop

Apache Hadoop architecture



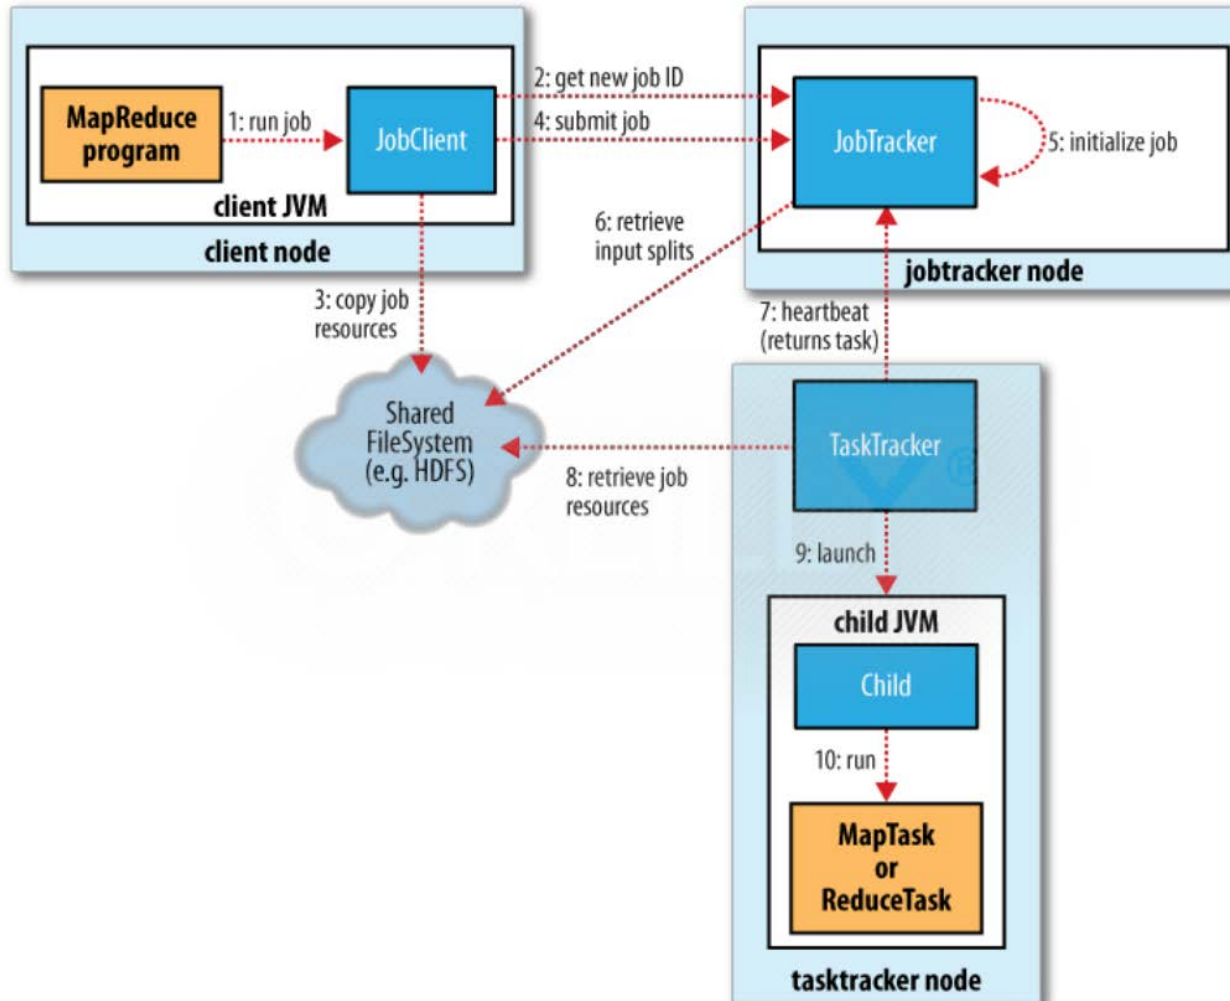
Apache Hadoop MapReduce data flow



CMSC 691 High Performance Distributed Systems

Apache Hadoop

Apache Hadoop MapReduce data flow



Apache Hadoop Installation

- Install software dependencies

```
$ sudo apt-get install maven git ssh rsync default-jdk default-jre openssh-server
```

- Download and install Eclipse Neon from <https://eclipse.org/downloads/>
- Download Apache Hadoop 2.7.3 from <http://hadoop.apache.org/releases.html> (binary) and extract in a folder e.g. /home/user/hadoop-2.7.3
- Edit line 25 of the file /home/user/hadoop-2.7.3/etc/hadoop/hadoop-env.sh

```
export JAVA_HOME=/usr/lib/jvm/default-java
```
- Edit your /home/user/.bashrc and append to the end:

```
export PATH=$PATH:/home/user/hadoop-2.7.3/bin
```
- Close the terminal and reopen it

CMSC 691 High Performance Distributed Systems

Apache Hadoop

Set up a new Eclipse project for Hadoop using Maven

- New -> Project -> Maven -> Maven Project
- Click on Next (default) two times
- Write a Group Id and Artifact Id, e.g. Hadoop and MapReduce
- Edit the pom.xml (right click, open with Text Editor) and add a new dependency within the dependencies section

```
<dependency>  
    <groupId>org.apache.hadoop</groupId>  
    <artifactId>hadoop-core</artifactId>  
    <version>1.2.1</version>  
</dependency>
```

- Delete the App.java and copy & paste the WordCount.java in the src/main/java as a baseline for your Hadoop MapReduce project

Compile and run a Eclipse project for Hadoop using Maven

- Right click on the project name -> Run As -> Maven install
- This will download dependencies, compile the code and package a jar e.g. in target/MapReduce-0.0.1-SNAPSHOT.jar
- Run Hadoop e.g. using the WordCount as:

```
hadoop jar target/MapReduce-0.0.1-SNAPSHOT.jar Hadoop.MapReduce.WordCount  
example/textfile.txt output
```

which states for:

```
hadoop jar jarfile.jar mainclass inputdata outputdata
```

- Analyze the files in the output
- Check out the online courses available in blackboard!

CMSC 691

High Performance Distributed Systems

Apache Hadoop

Dr. Alberto Cano
Assistant Professor
Department of Computer Science
acano@vcu.edu